CCIR · COMPUTE CREDIT INDEX RESEARCH

GPU-Backed Credit

Hypothetical Facility Covenant Framework

Training-to-Inference Migration · Covenant Trigger Analysis · March 2026

Hypothetical	GPU Co., LLC · $15M Senior Secured	**Collateral**	500 × NVIDIA H100 SXM
Term	36 months · Quarterly calculation dates	**Reference Rate**	CRI-H100 · ccir.io

This document sets out a hypothetical covenant framework for a GPU-backed senior secured credit facility, demonstrating how CRI-H100 and related CCIR reference rates can be cited by name in credit agreements. All borrower details, rates, and trigger levels are illustrative. This is not legal advice.

1. Background and Structuring Rationale

GPU-backed credit facilities have reached sufficient scale that the absence of a standardized, independent reference rate for compute rental prices creates material documentation and valuation risk. Without a named reference rate, every deal requires a bespoke valuation model — one that no counterparty can independently verify, and that creates interpretive friction at the worst possible moment: when a covenant is about to breach.

This framework addresses that gap by citing CRI-H100 — the interruptible spot floor published daily by Compute Credit Index Research — as the named reference rate across three covenant structures. The T2I Spread (Term-to-Interruptible) measures the haircut between going-concern cluster revenue and single-GPU spot redeployment value, serving as the empirical LGD input for the facility.

A second structural consideration motivates the workload composition covenant: the GPU collateral underpinning most H100-backed credit was underwritten as training hardware, but the market is actively migrating toward inference workloads as Blackwell-generation hardware displaces Hopper for large-model training. This migration changes the revenue stack the facility was underwritten against — and a covenant framework that does not account for it cannot adequately protect lender interests across a 36-month term.

1.1 The Training-to-Inference Transition

Three forces are compressing training demand on H100 collateral simultaneously:

- Model efficiency gains — each successive generation of frontier models requires less raw compute per parameter, reducing the aggregate H100 training hours required for equivalent capability.
- Inference-time compute shift — reasoning models front-load cost at inference rather than training, redistributing GPU hours toward interruptible spot workloads at lower per-hour revenue.
- Blackwell cascade — as B200 supply expands, H100 training premium erodes. When CRI-B200 falls within parity of CRI-H100, the H100's training use case is effectively over and the collateral is permanently repriced as inference-tier hardware.

The maintenance test, collateral coverage covenant, and workload composition trigger are designed to work together across this transition — each measuring a different dimension of the risk.

2. Defined Terms

The following definitions apply throughout this framework and are intended to be incorporated by reference into the applicable credit agreement.

"CRI-H100 Reference Rate" means, as of any date of determination, the daily interruptible spot floor rate for NVIDIA H100 SXM GPUs published by Compute Credit Index Research LLC (ccir.io) for the most recent Business Day preceding such date, expressed in United States dollars per GPU per hour, as computed pursuant to the CCIR Methodology (version then in effect).

"CRI-H100-CLS-MM" means the mid-market cluster rate for H100 SXM GPUs published by CCIR, representing the median of RunPod and Lambda 8-GPU posted prices, expressed in United States dollars per GPU per hour.

"T2I Spread" means, as of any date of determination, the percentage by which CRI-H100-CLS-MM exceeds the CRI-H100 Reference Rate, calculated as: (CRI-H100-CLS-MM minus CRI-H100) divided by CRI-H100-CLS-MM. A higher T2I Spread indicates greater distance between going-concern cluster revenue and spot redeployment value.

"CRI-B200 Reference Rate" means the daily interruptible spot floor rate for NVIDIA B200 SXM GPUs published by CCIR, expressed in United States dollars per GPU per hour.

"Generation Spread" means, as of any date of determination, the percentage by which the CRI-H100 Reference Rate exceeds the CRI-B200 Reference Rate, calculated as: (CRI-H100 minus CRI-B200) divided by CRI-H100. A Generation Spread at or below the Generation Spread Threshold indicates Blackwell pricing has converged with Hopper, signaling migration of training workloads to the next generation.

"GPU Rental Revenue Per Pledged GPU" means, for any Calculation Period, the aggregate GPU rental revenue received by the Borrower attributable to the Pledged GPUs, divided by the number of Pledged GPUs, divided by the number of hours in such Calculation Period, expressed as a dollar per GPU per hour figure.

"Pledged GPUs" means the 500 NVIDIA H100 SXM units pledged as collateral under the Security Agreement, as such number may be adjusted pursuant to the terms of the Credit Agreement.

"Calculation Date" means the last Business Day of each fiscal quarter of the Borrower.

"Training Workload Percentage" means, for any Calculation Period, the percentage of total GPU-hours attributable to the Pledged GPUs during such period that are consumed by model training workloads, as certified by the Borrower in the applicable Compliance Certificate.

"Rate Cap" means $2.80 per GPU per hour, being the ceiling applied to the CRI-H100 Reference Rate solely for purposes of computing the Maintenance Test Coverage Ratio, to protect a performing Borrower from technical breach arising from spot market appreciation unrelated to Borrower performance.

3. Covenant 1 — Maintenance Test

The maintenance test measures whether the Borrower's actual rental revenue per GPU remains adequate relative to what the spot market would bear for equivalent hardware. It is designed to catch deteriorating performance while protecting a borrower with locked contracts from breach driven solely by spot market appreciation.

3.1 Drafted Covenant Language

SECTION 7.01 — GPU REVENUE MAINTENANCE TEST — Financial Covenant

Borrower shall not permit the Maintenance Test Coverage Ratio, as of any Calculation Date, to be less than 2.50 to 1.00.

"Maintenance Test Coverage Ratio" means, as of any Calculation Date, the ratio of (a) GPU Rental Revenue Per Pledged GPU for the most recently completed Calculation Period to (b) the Capped CRI-H100 Reference Rate.

"Capped CRI-H100 Reference Rate" means, as of any Calculation Date, the lesser of (i) the CRI-H100 Reference Rate and (ii) the Rate Cap. For the avoidance of doubt, the Rate Cap is applied solely for purposes of this Section 7.01 and shall not apply to the calculation of Collateral Value under Section 7.02 or any other provision of this Agreement.

For purposes of this Section 7.01, if CRI-H100 is not published on a Calculation Date, the most recently published CRI-H100 Reference Rate within the prior five Business Days shall be used. If no publication is available within such period, the parties shall negotiate in good faith to agree on a substitute rate.

3.2 Structural Commentary

The Rate Cap serves two purposes. First, it prevents a borrower performing on locked multi-year contracts from breaching the maintenance test solely because the spot market has risen above their contracted rate — a circumstance that reflects market appreciation, not borrower deterioration. Second, it sets a ceiling on the lender's protection: if the rate cap is set at market underwriting, the lender retains covenant protection in any scenario where spot prices remain below underwritten levels.

The Rate Cap should be set at underwriting to reflect the spot rate at which the deal was sized. A Rate Cap materially above the underwritten rate provides the borrower with more protection than intended; a Rate Cap below the underwritten rate may cause immediate breach in normal market conditions.

Note the maintenance test can pass while the collateral coverage test deteriorates — a borrower with a locked contract is performing on revenue while the spot redeployment floor falls beneath them. Lenders should monitor both simultaneously.

4. Covenant 2 — Collateral Coverage / LTV

The collateral coverage covenant marks the pledged GPU fleet to market using CRI-H100 capitalized at an agreed yield rate — not the Borrower's contracted price. This ensures the lender's recovery analysis reflects what the hardware would fetch on redeployment into the spot market, independent of any single borrower's contracted revenue.

4.1 Drafted Covenant Language

SECTION 7.02 — COLLATERAL COVERAGE RATIO — Financial Covenant

Borrower shall not permit the Collateral Coverage Ratio, as of any Calculation Date, to be less than 1.25 to 1.00 (equivalent to a maximum loan-to-value of 80%).

"Collateral Coverage Ratio" means, as of any Calculation Date, the ratio of (a) Collateral Value to (b) the outstanding principal amount of the Loans.

"Collateral Value" means, as of any Calculation Date, the product of (i) the number of Pledged GPUs, multiplied by (ii) the CRI-H100 Reference Rate (without application of the Rate Cap), multiplied by (iii) 8,760 (annual GPU hours), multiplied by (iv) the Utilization Assumption, divided by (v) the Capitalization Rate.

"Utilization Assumption" means 0.72 (72%), representing a conservative estimate of productive utilization in an interruptible spot deployment scenario. The Administrative Agent may adjust the Utilization Assumption downward (but not upward) by written notice if spot market utilization data from CCIR methodology publications supports a lower figure.

"Capitalization Rate" means 0.28 (28%), representing the agreed yield rate applied to annualized spot GPU rental income to derive a capitalized collateral value. The Capitalization Rate shall be reviewed annually and may be adjusted by mutual agreement of the parties.

Upon any Collateral Coverage Ratio Breach, Borrower shall, within 30 days of notice from the Administrative Agent, either (a) prepay Loans in an amount sufficient to restore compliance, or (b) pledge additional Pledged GPUs or other collateral acceptable to the Administrative Agent in its reasonable discretion.

4.2 Structural Commentary

The Collateral Value formula is a cap rate methodology analogous to those used in real estate and infrastructure ABS: annualized rental income divided by a yield rate. The key structural distinction from traditional asset lending is that the income stream — and therefore the asset value — is marked to a live reference rate rather than a fixed contracted price. This is both the primary benefit and the primary source of mark-to-market volatility in the covenant.

The Utilization Assumption of 72% reflects the discount between theoretical maximum GPU hours and realistic spot market deployment — accounting for scheduling gaps, platform downtime, and the interruptible nature of spot contracts. Lenders should consider whether to include a contractual floor utilization (e.g., if the borrower can demonstrate contracted utilization above 85%, the Utilization Assumption for that tranche may be adjusted upward).

The Capitalization Rate is the most negotiated variable in this structure. A lower cap rate (e.g., 20%) produces a higher Collateral Value and more favorable LTV; a higher cap rate (e.g., 35%) is more conservative. As the GPU rental market matures and transaction comps accumulate, market consensus cap rates will likely emerge — a function CCIR's published rate history is designed to support.

5. Covenant 3 — Workload Composition Trigger

The workload composition covenant addresses the training-to-inference migration risk directly. Because H100 training revenue per GPU is materially higher than inference revenue per GPU, a borrower who quietly migrates their GPU fleet from training to inference workloads can simultaneously satisfy a blunt revenue-per-GPU test while the revenue stack the facility was underwritten against has fundamentally changed.

5.1 Drafted Covenant Language

SECTION 7.03 — WORKLOAD COMPOSITION COVENANT — Operational Covenant

Borrower shall not permit the Training Workload Percentage, as of any two consecutive Calculation Dates, to be less than 50%.

If the Training Workload Percentage falls below 50% for any single Calculation Period (a "Workload Migration Event"), Borrower shall deliver written notice to the Administrative Agent within 10 Business Days, which notice shall include: (i) a description of the workloads consuming the redeployed GPU capacity; (ii) the contracted rate and term for such workloads; and (iii) a revised revenue forecast for the next two Calculation Periods.

If the Training Workload Percentage falls below 50% for two consecutive Calculation Periods (a "Workload Migration Breach"), the Collateral Value for purposes of Section 7.02 shall be computed using a Blended Utilization Rate rather than the standard Utilization Assumption.

"Blended Utilization Rate" means the weighted average of (a) the Training Utilization Assumption of 0.72, weighted by the Training Workload Percentage, and (b) the Inference Utilization Assumption of 0.65, weighted by one minus the Training Workload Percentage. The lower blended rate reflects the lower per-hour revenue and higher interruptibility of inference workloads relative to training.

For the avoidance of doubt, a Workload Migration Breach is not itself an Event of Default; it operates solely to adjust the Collateral Value calculation. An Event of Default arising from a subsequent Collateral Coverage Ratio Breach shall be governed by Section 8.01.

5.2 Structural Commentary

The workload composition covenant is an information covenant as much as a financial covenant — the Workload Migration Event notice requirement is designed to surface the transition before it becomes a breach, giving the lender visibility into the revenue stack change in real time rather than at the point of a collateral coverage breach.

The 50% training threshold is a starting point for negotiation and should be calibrated to the borrower's actual revenue model at underwriting. A borrower with a pure training business warrants a higher threshold (e.g., 65%); a borrower with a mixed training/inference book at origination may negotiate a lower one.

Critically, the covenant does not penalize workload migration per se — inference demand going up is not a credit event. It adjusts the collateral mark to reflect the changed revenue profile. A borrower who migrates to inference but maintains high utilization and adequate contracted revenue may satisfy the collateral coverage test even under the Blended Utilization Rate.

6. Generation Spread Covenant — Blackwell Cascade Protection

The generation spread covenant addresses the forward-looking risk that Blackwell-generation hardware (B200) reaches price parity with Hopper (H100), effectively eliminating the training premium that justifies H100 collateral valuations. This is not a near-term risk in current market conditions, but over a 36-month facility term it is a material tail scenario.

6.1 Drafted Covenant Language

SECTION 7.04 — GENERATION SPREAD COVENANT — Structural Protection

If, as of any Calculation Date, the Generation Spread is less than the Generation Spread Threshold (15%), Borrower shall deliver written notice to the Administrative Agent within 10 Business Days (a "Generation Convergence Notice").

If, as of any two consecutive Calculation Dates, the Generation Spread is less than the Generation Spread Threshold (a "Generation Convergence Event"), the following adjustments shall take effect automatically on the next Calculation Date:

(a) Accelerated Amortization: Borrower shall prepay Loans at a rate of not less than 10% of original principal per Calculation Period until the Generation Spread returns to or above the Generation Spread Threshold for two consecutive Calculation Periods; and

(b) Collateral Remarking: the Capitalization Rate in Section 7.02 shall increase from 28% to 35%, reflecting the higher yield required to compensate for collateral that no longer commands a training-tier revenue premium.

The Generation Spread Threshold of 15% represents the level below which B200 spot pricing has effectively arbitraged the H100 training premium, and the market is treating H100 hardware as inference-tier rather than training-tier collateral. The Administrative Agent and the Borrower shall review the Generation Spread Threshold annually and may adjust it by mutual written agreement.

6.2 Structural Commentary

The generation spread covenant is a forward-looking protection that most GPU-backed credit documents currently lack entirely. Its primary function is to prevent a lender from holding a 36-month facility against collateral whose revenue tier has structurally degraded without any contractual mechanism to respond.

The accelerated amortization provision is preferable to an immediate Event of Default because generation convergence is a market-driven event, not a borrower performance failure. A performing borrower whose collateral is being commoditized by market forces should be given time to amortize the facility down, not immediately accelerated.

The Capitalization Rate step-up from 28% to 35% adjusts the collateral mark to reflect a higher required yield on inference-tier hardware — consistent with how infrastructure investors price assets with shorter effective economic lives or higher obsolescence risk.

7. Scenario Analysis — Covenant Behavior Under Market Stress

The following table illustrates how each covenant responds across five market scenarios. The hypothetical facility has a Borrower with 500 H100 SXM GPUs, contracted revenue of $2.80/GPU/hr on a 3-year locked contract, and $15M in outstanding principal. Maintenance Test minimum: 2.50×. Collateral Coverage minimum: 1.25× (80% LTV). Workload Trigger: training below 50% for two consecutive periods.

Scenario	CRI-H100 Spot	CLS-MM	T2I Spread	Maintenance Test	Collateral Coverage	Workload Trigger
Baseline	$2.05	$3.40	39.7%	PASS — 2.93×	PASS — 68%	PASS — 80% training
Spot Decline (inference cascade)	$1.45	$3.20	54.7%	PASS — contracted revenue holds	WARN — collateral marks down	WATCH — training at 65%
Spot Rise (demand surge)	$3.10	$4.50	31.1%	WARN — Rate Cap engaged; test uses $2.80 ceiling	PASS — collateral marks up	PASS — 75% training
Training Collapse (workload migration)	$1.20	$2.80	57.1%	PASS — locked contracts hold	BREACH — LTV > 85%	BREACH — training at 35%
Blackwell Cascade (B200 ≈ H100)	$1.05	$2.40	56.3%	PASS — legacy contracts intact	BREACH — generation spread triggers prepayment	BREACH — training demand migrating to B200

Key observations from the scenario analysis:

- Locked contracts decouple the maintenance test from spot price movements — a falling floor helps the borrower on covenant 1 while simultaneously hurting on covenant 2. Lenders must monitor both.
- The Rate Cap matters most in a rising market scenario — without it, a Borrower earning $2.80/hr against a $3.10 spot floor would show a 0.90× coverage ratio and breach, despite performing exactly on contract.
- Training Collapse and Blackwell Cascade scenarios both fire covenant 3 and covenant 2 simultaneously — by design, since workload migration and generation convergence are both ultimately reflected in the spot floor.
- No single scenario causes all three covenants to breach simultaneously — the framework is designed to give graduated signals, not a binary default trigger.

8. CCIR Covenant Structuring Consultation

The covenant framework set out in this paper represents one configuration of CRI-H100 and related CCIR indices in a GPU-backed credit facility. Deal-specific structuring requires calibration of covenant thresholds, Rate Cap levels, Capitalization Rates, and trigger mechanisms to the actual borrower, collateral pool, and credit profile.

CCIR offers bespoke covenant structuring consultation for lenders, borrowers, and their counsel working on GPU-backed credit transactions. Consultation services include:

- Threshold calibration — setting Rate Cap, coverage ratio minimums, workload trigger percentages, and Generation Spread Threshold to the specific deal economics at underwriting.
- Collateral pool analysis — reviewing the GPU mix, contract structure, and workload profile to identify which CCIR indices are relevant and how they should be referenced in the covenant package.
- Covenant language review — working with transaction counsel to ensure CRI index citations are precise, methodology references are current, and fallback provisions are adequate.
- Ongoing rate monitoring — providing lenders with Calculation Date rate packages showing CRI-H100, CLS-MM, CLS-ENT, T2I Spread, and Generation Spread in a format designed for compliance certificate review.

Engage CCIR for your next GPU-backed credit transaction

Independent rate administration · Covenant structuring consultation · Ongoing compliance support

research@ccir.io
ccir.io

Disclaimer

This document is a working paper published by Compute Credit Index Research LLC for informational purposes. It is not legal advice and does not constitute a commitment to provide financing or any other financial service. All borrower details, rates, and covenant thresholds are hypothetical and illustrative. Transaction parties should seek independent legal counsel when structuring credit agreements. CCIR holds no financial interest in the GPU compute market it measures.